

```
import numpy as np
import torch
```

# Gradient Descent Practice

```
inputs = np.array([[73, 67, 43],
                   [91, 88, 64],
                   [87, 134, 58],
                   [102, 43, 37],
                   [69, 96, 70]], dtype='float32')
```

```
targets = np.array([[56, 70],
                    [81, 101],
                    [119, 133],
                    [22, 37],
                    [103, 119]], dtype='float32')
```

```
inputs = torch.from_numpy(inputs)
targets = torch.from_numpy(targets)
```

```
# Equation Looks like:
# apple = w11*x + w12*y + w13*z + b1
# orange = w21*x + w22*y + w23*z + b2

# Weight Initialization:
w = torch.randn(size=(2, 3), requires_grad=True) # 2x3
# Biases Initialization:
b = torch.randn(size=(2, ), requires_grad=True)

display(w)
display(b)
```

```
tensor([[ -1.3902, -0.1076, -1.2027],
        [ 0.1122, -0.8120,  0.7782]], requires_grad=True)
tensor([-0.2583,  0.3550], requires_grad=True)
```

```
# Model
def model(inputs):
    return inputs @ w.t() + b
```

```
# Loss
def loss_(predictions, truths):
    diff = predictions - truths

    return torch.sum(diff * diff) / diff.numel()
```

```
# Full Training
epochs = 100
learning_rate = 3e-5
```

```
for epoch in range(epochs):
    print(f"Epoch Number: {epoch}: ")
    predictions = model(inputs)
    loss = loss_(predictions, targets)
    print(f"\tLoss: ", loss.detach().numpy())
    loss.backward()
    with torch.no_grad():
        factor = w*learning_rate
        factor_2 = b*learning_rate
        w.sub_(factor)
        b.sub_(factor_2)
```

```
Epoch Number: 0:
    Loss:  44361.797
Epoch Number: 1:
    Loss:  44360.145
Epoch Number: 2:
    Loss:  44358.496
Epoch Number: 3:
    Loss:  44356.844
Epoch Number: 4:
    Loss:  44355.19
Epoch Number: 5:
    Loss:  44353.543
Epoch Number: 6:
    Loss:  44351.883
Epoch Number: 7:
    Loss:  44350.24
Epoch Number: 8:
    Loss:  44348.586
Epoch Number: 9:
    Loss:  44346.938
Epoch Number: 10:
    Loss:  44345.29
Epoch Number: 11:
    Loss:  44343.633
Epoch Number: 12:
    Loss:  44341.98
Epoch Number: 13:
    Loss:  44340.332
Epoch Number: 14:
```

Loss: 44338.68  
Epoch Number: 15:  
Loss: 44337.03  
Epoch Number: 16:  
Loss: 44335.38  
Epoch Number: 17:  
Loss: 44333.734  
Epoch Number: 18:  
Loss: 44332.08  
Epoch Number: 19:  
Loss: 44330.43  
Epoch Number: 20:  
Loss: 44328.773  
Epoch Number: 21:  
Loss: 44327.13  
Epoch Number: 22:  
Loss: 44325.48  
Epoch Number: 23:  
Loss: 44323.83  
Epoch Number: 24:  
Loss: 44322.18  
Epoch Number: 25:  
Loss: 44320.54  
Epoch Number: 26:  
Loss: 44318.883  
Epoch Number: 27:  
Loss: 44317.24  
Epoch Number: 28:  
Loss: 44315.586  
Epoch Number: 29:  
Loss: 44313.945  
Epoch Number: 30:  
Loss: 44312.29  
Epoch Number: 31:  
Loss: 44310.65  
Epoch Number: 32:  
Loss: 44308.992  
Epoch Number: 33:  
Loss: 44307.344  
Epoch Number: 34:  
Loss: 44305.695  
Epoch Number: 35:  
Loss: 44304.055

Epoch Number: 36:  
Loss: 44302.406

Epoch Number: 37:  
Loss: 44300.766

Epoch Number: 38:  
Loss: 44299.113

Epoch Number: 39:  
Loss: 44297.47

Epoch Number: 40:  
Loss: 44295.824

Epoch Number: 41:  
Loss: 44294.18

Epoch Number: 42:  
Loss: 44292.53

Epoch Number: 43:  
Loss: 44290.89

Epoch Number: 44:  
Loss: 44289.242

Epoch Number: 45:  
Loss: 44287.6

Epoch Number: 46:  
Loss: 44285.95

Epoch Number: 47:  
Loss: 44284.31

Epoch Number: 48:  
Loss: 44282.664

Epoch Number: 49:  
Loss: 44281.023

Epoch Number: 50:  
Loss: 44279.375

Epoch Number: 51:  
Loss: 44277.727

Epoch Number: 52:  
Loss: 44276.08

Epoch Number: 53:  
Loss: 44274.44

Epoch Number: 54:  
Loss: 44272.793

Epoch Number: 55:  
Loss: 44271.15

Epoch Number: 56:  
Loss: 44269.508

Epoch Number: 57:

Loss: 44267.86  
Epoch Number: 58:  
Loss: 44266.22  
Epoch Number: 59:  
Loss: 44264.57  
Epoch Number: 60:  
Loss: 44262.926  
Epoch Number: 61:  
Loss: 44261.285  
Epoch Number: 62:  
Loss: 44259.633  
Epoch Number: 63:  
Loss: 44257.992  
Epoch Number: 64:  
Loss: 44256.35  
Epoch Number: 65:  
Loss: 44254.703  
Epoch Number: 66:  
Loss: 44253.055  
Epoch Number: 67:  
Loss: 44251.414  
Epoch Number: 68:  
Loss: 44249.77  
Epoch Number: 69:  
Loss: 44248.13  
Epoch Number: 70:  
Loss: 44246.484  
Epoch Number: 71:  
Loss: 44244.84  
Epoch Number: 72:  
Loss: 44243.195  
Epoch Number: 73:  
Loss: 44241.55  
Epoch Number: 74:  
Loss: 44239.902  
Epoch Number: 75:  
Loss: 44238.26  
Epoch Number: 76:  
Loss: 44236.617  
Epoch Number: 77:  
Loss: 44234.973  
Epoch Number: 78:  
Loss: 44233.33

Epoch Number: 79:  
Loss: 44231.684

Epoch Number: 80:  
Loss: 44230.04

Epoch Number: 81:  
Loss: 44228.4

Epoch Number: 82:  
Loss: 44226.758

Epoch Number: 83:  
Loss: 44225.11

Epoch Number: 84:  
Loss: 44223.473

Epoch Number: 85:  
Loss: 44221.83

Epoch Number: 86:  
Loss: 44220.18

Epoch Number: 87:  
Loss: 44218.54

Epoch Number: 88:  
Loss: 44216.895

Epoch Number: 89:  
Loss: 44215.254

Epoch Number: 90:  
Loss: 44213.605

Epoch Number: 91:  
Loss: 44211.96

Epoch Number: 92:  
Loss: 44210.32

Epoch Number: 93:  
Loss: 44208.68

Epoch Number: 94:  
Loss: 44207.035

Epoch Number: 95:  
Loss: 44205.395

Epoch Number: 96:  
Loss: 44203.742

Epoch Number: 97:  
Loss: 44202.1

Epoch Number: 98:  
Loss: 44200.46

Epoch Number: 99:  
Loss: 44198.816

```
final_predictions = model(inputs)
display(final_predictions)
display(targets)
```

```
tensor([[ -160.1831,  -12.3593],
        [-212.5647,  -11.0532],
        [-204.7594,  -53.3977],
        [-190.6086,   5.6595],
        [-190.1254, -15.3354]], grad_fn=<AddBackward0>)
tensor([[ 56.,  70.],
        [ 81., 101.],
        [119., 133.],
        [ 22.,  37.],
        [103., 119.]])
```

## Questions for Review

Try answering the following questions to test your understanding of the topics covered in this notebook:

1. What is a linear regression model? Give an example of a problem formulated as a linear regression model.
  - In linear regression, we assume a relation between inputs and the targets and then we weigh the inputs based on their importance by calculating their gradients.
2. What are input and target variables in a dataset? Give an example.
  - Input: The input to the model. Can be called features too. Target: Truth values.
3. What are weights and biases in a linear regression model?
  - Weights: Importance of a specific feature/input. Bias: A threshold / cutoff for the targets.
4. How do you represent tabular data using PyTorch tensors?
  - `torch.tensor`
5. Why do we create separate matrices for inputs and targets while training a linear regression model?
  - Because we have to calculate the loss.
6. How do you determine the shape of the weights matrix & bias vector given some training data?
  - Weight matrix: (number of targets, number of features). Bias Vector: (number of targets).
7. How do you create randomly initialized weights & biases with a given shape?
  - `torch.randn(size=())`
8. How is a linear regression model implemented using matrix operations? Explain with an example. Done.
9. How do you generate predictions using a linear regression model?
  - `model(inputs)`
10. Why are the predictions of a randomly initialized model different from the actual targets?

- Because the weights aren't right.

11. What is a loss function? What does the term “loss” signify?

- The error in predictions of the model.

12. What is mean squared error?

- $\text{error}^2 / \text{num\_of\_inputs}$

13. Write a function to calculate mean squared using model predictions and actual targets.

- Done

14. What happens when you invoke the `.backward` function on the result of the mean squared error loss function?

- Calculates gradient

15. Why is the derivative of the loss w.r.t. the weights matrix itself a matrix? What do its elements represent?

- How much do we have to penalize the weights.

16. How is the derivate of the loss w.r.t. a weight element useful for reducing the loss? Explain with an example.

- We have to lower or increase the weight so that it can approach the minimum.

17. Suppose the derivative of the loss w.r.t. a weight element is positive. Should you increase or decrease the element's value slightly to get a lower loss?

- Decrease

18. Suppose the derivative of the loss w.r.t. a weight element is negative. Should you increase or decrease the element's value slightly to get a lower loss?

- Increase

19. How do you update the weights and biases of a model using their respective gradients to reduce the loss slightly?

- Multiply the weight derivative by learning rate.

20. What is the gradient descent optimization algorithm? Why is it called “gradient descent”?

- We are descending towards local minimum.

21. Why do you subtract a “small quantity” proportional to the gradient from the weights & biases, not the actual gradient itself?

- Learning rate.

22. What is learning rate? Why is it important?

- A factor with which the weights are to be managed. Important to reach minima.

23. What is `torch.no_grad`?

- Explicitly tell pytorch to not calculate gradients because they are being done in the context explicitly.



24. Why do you reset gradients to zero after updating weights and biases?

- Don't add up.

25. What are the steps involved in training a linear regression model using gradient descent?

- 1. Predictions, 2. Loss, 3. Gradients, 4. Update weights.

26. What is an epoch?

- Cycle.

27. What is the benefit of training a model for multiple epochs?

- Slowly slowly descending towards local minimum.

28. How do you make predictions using a trained model?

- `model(inputs)`

29. What should you do if your model's loss doesn't decrease while training? Hint: learning rate.

- Lower learning rate.

30. What is `torch.nn`?

- neural network module of pytorch

31. What is the purpose of the `TensorDataset` class in PyTorch? Give an example.

- Convert to train test batch.

32. What is a data loader in PyTorch? Give an example.

- For batching of data.

33. How do you use a data loader to retrieve batches of data?

- Updating of model with batches.

34. What are the benefits of shuffling the training data before creating batches?

- Randomize the targets variable.

35. What is the benefit of training in small batches instead of training with the entire dataset?

- To learn the weights sequentially.

36. What is the purpose of the `nn.Linear` class in PyTorch? Give an example.

- Linear module.

37. How do you see the weights and biases of a `nn.Linear` model?

- `model.weight`, `model.bias`

38. What is the purpose of the `torch.nn.functional` module?

- For loss functions

39. How do you compute mean squared error loss using a PyTorch built-in function?

- `F.mse_loss`

40. What is an optimizer in PyTorch?

- helps achieve gradient descent quicker.

41. What is `torch.optim.SGD`? What does SGD stand for?

- stochastic gradient descent.

42. What are the inputs to a PyTorch optimizer?

- model's parameters and learning rate.

43. Give an example of creating an optimizer for training a linear regression model.

- `torch.optim.SGD(model.parameters(), lr=1e-5)`

44. Write a function to train a `nn.Linear` model in batches using gradient descent.

- Done

45. How do you use a linear regression model to make predictions on previously unseen data?

- `model(unseen_input)`